

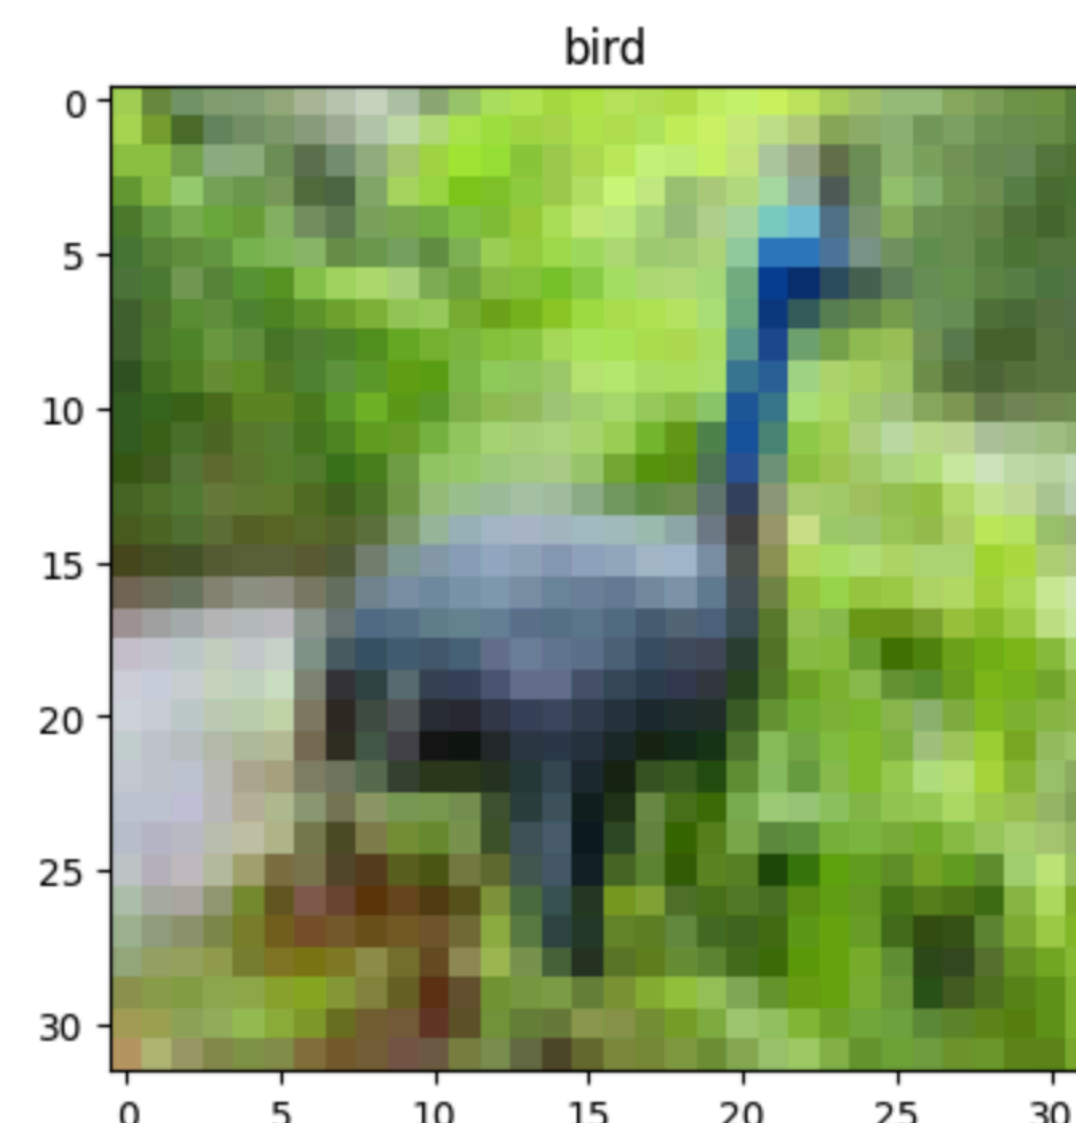
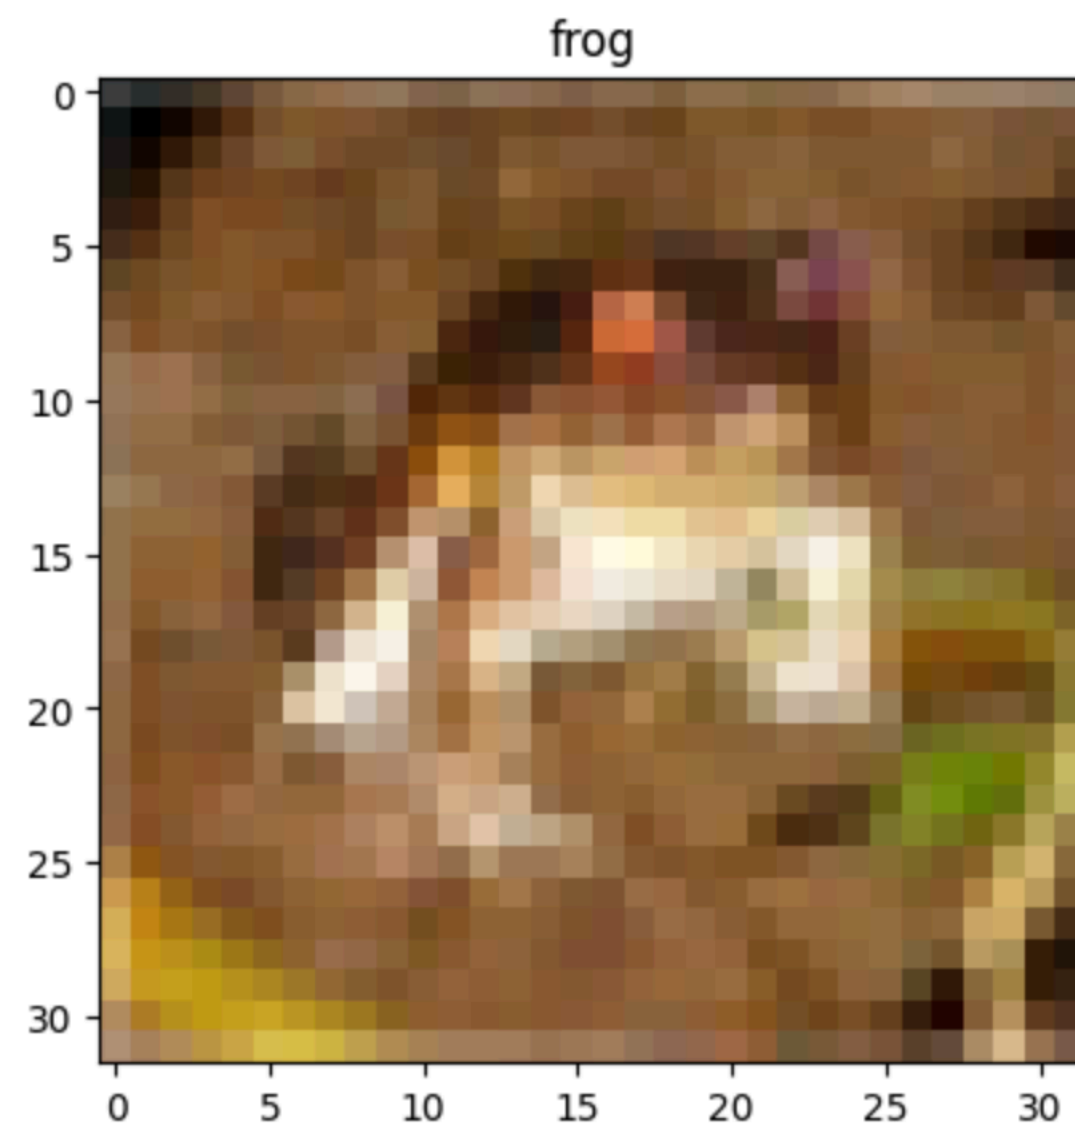
# Computer Vision Project

Build and train a model with CIFAR-10 dataset

# Project Goal

## Image classification using deep learning and transfer learning

- Image classification with CIFAR-10
- Focus on deep learning workflow
- Training and fine-tuning a CNN model



# Model Architecture

## Combining ResNet50 with a custom classification head

- Base model: ResNet50 (retrained on ImageNet, top removed)
- Added layers: GlobalAveragePooling + Dense (512 → 128)
- Output: Softmax layer with 10 classes
- Used ReLu and Dropout for regularization

```
# Load base model
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
base_model.trainable = False # Freeze base model

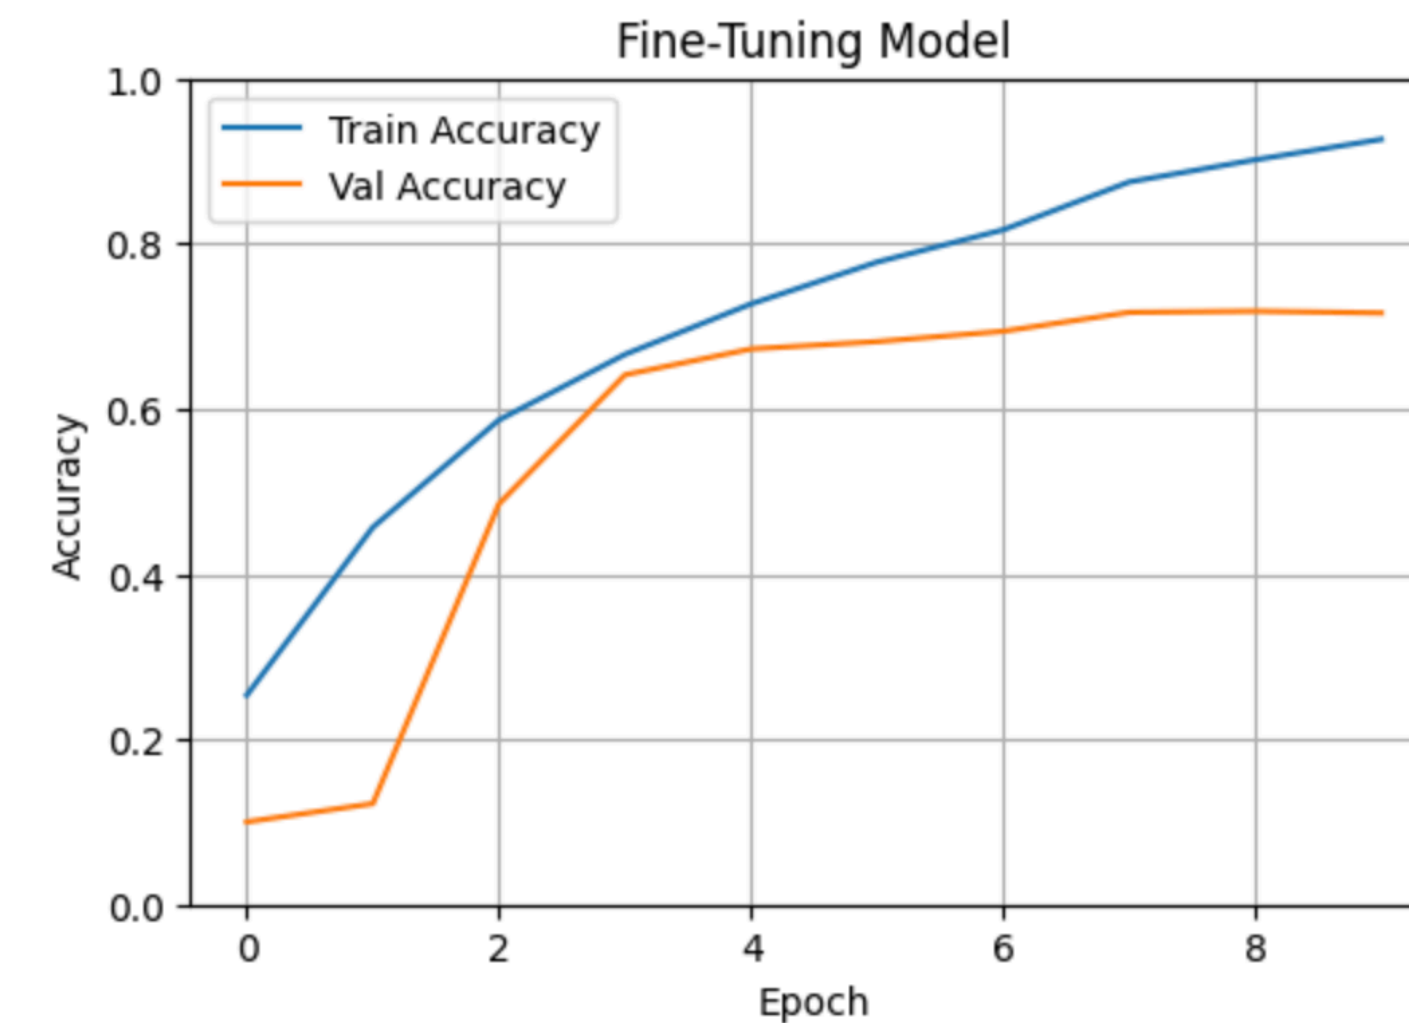
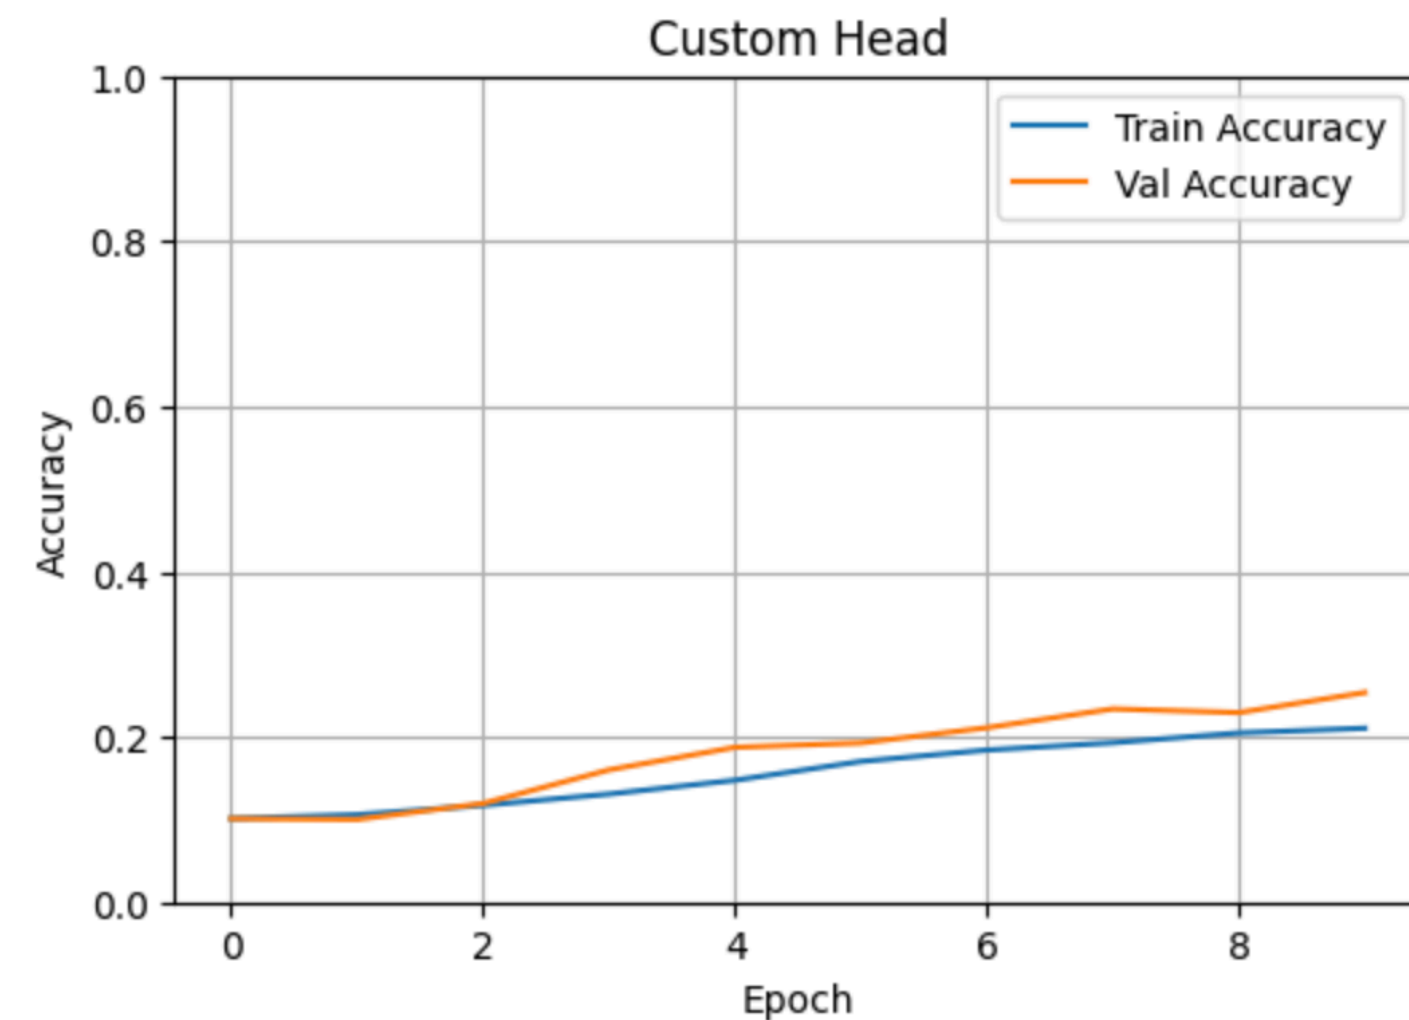
# Add custom classification head
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu')(x)
x = Dropout(0.3)(x)
x = Dense(128, activation='relu')(x)
x = Dropout(0.2)(x)
predictions = Dense(10, activation='softmax')(x)

# Combine base model and head
model = Model(inputs=base_model.input, outputs=predictions)
```

# Training Strategy

## From initial training to fine-tuning with callbacks and augmentation

- Phase 1: Train classification head with frozen base model
- Phase 2: Fine-tune entire model with low learning rate
- Used Adam optimiser, learning rate scheduler & early stopping
- Optional: image augmentation (rotation, shift, flips)



# Results & Evaluation

## Visualizing performance and assessing predictions

- Final validation accuracy: ~ 73%
- Confusion matrix shows where the model struggles
- Sample predictions show incorrect classifications

Confusion Matrix


airplane	701	20	73	7	23	3	14	15	99	45
automobile	22	817	8	14	3	3	7	9	37	80
bird	53	5	629	59	112	47	60	25	8	2
cat	19	9	55	555	76	119	79	49	25	14
deer	25	2	69	48	674	25	54	86	14	3
dog	10	1	50	198	57	563	30	76	10	5
frog	10	11	58	50	41	13	805	4	3	5
horse	9	2	40	41	54	31	12	789	2	20
ship	46	37	9	13	10	1	2	5	845	32
truck	24	90	9	19	9	2	4	15	39	789
	airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck

Predicted Label


True: automobile  
Pred: truck



True: bird  
Pred: cat




True: dog  
Pred: cat



True: deer  
Pred: airplane



True: dog  
Pred: deer



True: horse  
Pred: deer



# Conclusion & Next Steps

## Reflections on building and future improvements

- I learned to build, train and evaluate a CNN with transfer learning
- I explored model optimization with different layers and callbacks
- For further improvements:
  - Try other architectures (e.g., EfficientNet)
  - Use more data or advanced augmentation
  - Optimize hyperparameters

# Thank You!

## Questions, feedback, ideas?

- Thank you for your attention.
- Feel free to ask questions or share feedback.
- I'm happy to discuss challenges or improvements in image classification.
- Let me know if you want access to to the notebook or code.
- Connect with me on [LinkedIn](#)